

Week 4: Graph Data Structures

Week 4

1/105

Things to Note ...

- Online help sessions: check the course homepage for up-to-date timetable
⇒ *start weekly assessments *early* - tutors can help you better if you do*
- Mid term - Moodle quiz, Week 6 Thursday 2pm-3pm, details to follow

In This Lecture ...

- Graph data structures (Slides, [S] Ch. 17.1-17.5)
- Graph traversal (Slides, [S] Ch. 17.7, 18.1-18.3, 18.7)

Coming Up ...

- Algorithms on graphs ([S] Ch. 19-22)

Nerds You Should Know

2/105

The next in a series on famous computer scientists ...



What he invented affects your life every single day ...

... Nerds You Should Know

3/105

Sir Tim Berners-Lee

- Oxford CS Graduate (1976)
- Software engineer at CERN
- Founder/Director of W3C at MIT (1994)

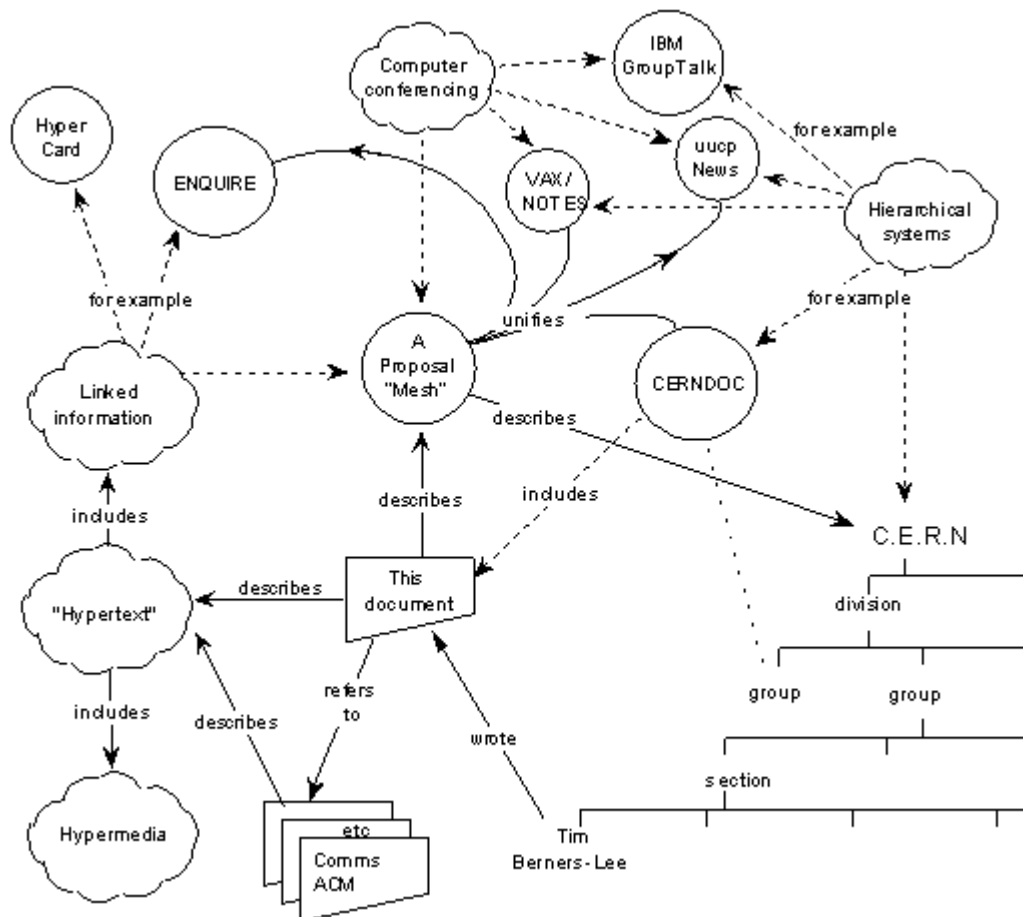


- Inventing the Web ...
 - distributed hypertext
 - linking heterogeneous documents
 - universal naming scheme (URL)
 - transfer protocol (http)
 - later apologised for initial pair of slashes ('// ') in a web address
 - also thinks he should have defined web addresses the other way round (au.edu.unsw.cse)
- Winner of the Turing Award in 2016

... Nerds You Should Know

4/105

Tim Berners-Lee's original diagram of the "Web"



(from his proposal document, 1989)

Graph Definitions

Graphs

6/105

Many applications require

- a collection of *items* (i.e. a set)
- *relationships/connections* between items

Examples:

- maps: items are cities, connections are roads
- web: items are pages, connections are hyperlinks

Collection types you're familiar with

- arrays, lists ... linear sequence of items

Graphs are more general ... allow arbitrary connections


... Graphs

7/105

A graph $G = (V, E)$

- V is a set of *vertices*
- E is a set of *edges* (subset of $V \times V$)

Example:



... Graphs

8/105

A real example: Australian road distances


Distance	Adelaide	Brisbane	Canberra	Darwin	Melbourne	Perth	Sydney
Adelaide	-	2055	-	3051	732	2716	-
Brisbane	2055	-	-	3429	1671	-	982
Canberra	-	-	-	-	658	-	309
Darwin	3051	3429	-	-	-	4049	-
Melbourne	732	1671	658	-	-	-	873
Perth	2716	-	-	4049	-	-	-
Sydney	-	982	309	-	873	-	-

Notes: vertices are cities, edges are distance between cities, symmetric

... Graphs

9/105

Alternative representation of above:



... Graphs

10/105

Questions we might ask about a graph:

- is there a way to get from item A to item B?
- what is the best way to get from A to B?
- which items are connected?

Graph algorithms are generally more complex than tree/list ones:

- no implicit order of items (cf. linked lists)
- graphs may contain cycles
- concrete representation (i.e., a data structure) is less obvious
- algorithm complexity depends on connection complexity

Properties of Graphs

11/105

For now, consider a graph with no (a,a) and no direction, no weight

Terminology: $|V|$ and $|E|$ (cardinality) normally written just as V and E .

A graph with V vertices has at most $V(V-1)/2$ edges.

The ratio $E:V$ can vary considerably.

- if E is closer to V^2 , the graph is *dense*
- if E is closer to V , the graph is *sparse*
 - Example: web pages and hyperlinks

Knowing whether a graph is sparse or dense is important

- may affect choice of data structures to represent graph
- may affect choice of algorithms to process graph

Exercise #1: Number of Edges

12/105

The edges in a graph represent pairs of connected vertices. A graph with V has V^2 such pairs.

Consider $V = \{1,2,3,4,5\}$ with all possible pairs:

$$E = \{ (1,1), (1,2), (1,3), (1,4), (1,5), (2,1), (2,2), \dots, (4,5), (5,5) \}$$

Why do we say that the maximum #edges is $V(V-1)/2$?

... because

- (v,w) and (w,v) denote the same edge (in an undirected graph)
- we do not consider loops (v,v)

Graph Terminology

14/105

For an edge e that connects vertices v and w

- v and w are *adjacent* (neighbours)
- e is *incident* on both v and w

Degree of a vertex v

- number of edges incident on v

Synonyms:

- vertex = node, edge = arc = link (Note: some people use arc for *directed* edges)

... Graph Terminology

15/105

Path: a sequence of vertices where

- each vertex has an edge to its predecessor

Simple path: a path where


- all vertices and edges are different

Cycle: a path

- that is simple except last vertex = first vertex

Length of path or cycle:

- #edges



... Graph Terminology

16/105

Connected graph

- there is a *path* from each vertex to every other vertex
- if a graph is not connected, it has ≥ 2 *connected components*

Complete graph K_V

- there is an *edge* from each vertex to every other vertex
- in a complete graph, $E = V(V-1)/2$



... Graph Terminology


17/105

Tree: connected (sub)graph with no cycles

Spanning tree: tree containing all vertices

Clique: complete subgraph

Consider the following single graph:



This graph has 26 vertices, 33 edges, and 4 connected components

Note: The entire graph has no spanning tree; what is shown in green is a spanning tree of the third connected component

... Graph Terminology

18/105


Undirected graph

- $edge(u,v) = edge(v,u)$, no self-loops (i.e. no $edge(v,v)$)

Directed graph

- $edge(u,v) \neq edge(v,u)$, can have self-loops (i.e. $edge(v,v)$)

Examples:



... Graph Terminology

19/105

Other types of graphs ...

Weighted graph

- each edge has an associated value (weight)
- e.g. road map (weights on edges are distances between cities)

Multi-graph

- allow multiple edges between two vertices
- e.g. function call graph ($f()$ calls $g()$ in several places)

Graph Data Structures


Graph Representations

21/105

Defining graphs:

- need some way of identifying vertices
- could give diagram showing edges and vertices
- could give a list of edges

E.g. four representations of the same graph:



... Graph Representations

22/105

We will discuss three different graph data structures:

1. Array of edges
2. Adjacency matrix
3. Adjacency list

Array-of-edges Representation

23/105

Edges are represented as an array of Edge values (= pairs of vertices)

- space efficient representation
- adding and deleting edges is slightly complex
- undirected: order of vertices in an Edge doesn't matter
- directed: order of vertices in an Edge encodes direction



For simplicity, we always assume vertices to be numbered $0 \dots V-1$

... Array-of-edges Representation

24/105

Graph initialisation (non-directed graph)

```

newGraph(V):
|   Input   number of nodes V
|   Output new empty graph
|
|   g.nV = V    // #vertices (numbered 0..V-1)
|   g.nE = 0    // #edges
|   allocate enough memory for g.edges[]
|   return g

```

How much is enough? ... No more than $V(V-1)/2$... Much less in practice (sparse graph)

... Array-of-edges Representation

25/105

Edge insertion

```

insertEdge(g, (v,w)):
|   Input graph g, edge (v,w) // assumption: (v,w) not in g
|
|   g.edges[g.nE]=(v,w)
|   g.nE=g.nE+1

```

... Array-of-edges Representation

26/105

Edge removal

```

removeEdge(g, (v,w)):
|   Input graph g, edge (v,w) // assumption: (v,w) in g
|
|   i=0
|   while (v,w)≠g.edges[i] do
|       i=i+1
|   end while
|   g.edges[i]=g.edges[g.nE-1] // replace (v,w) by last edge in array
|   g.nE=g.nE-1

```

Cost Analysis

27/105

Storage cost: $O(E)$ (big enough for # of edges)

Cost of operations:

- initialisation: $O(1)$
- insert edge: $O(1)$ (assuming edge array has space, edge does not exist in g)
- find/delete edge: $O(E)$ (need to find edge in edge array)

If array is full on insert

- allocate space for a bigger array, copy edges across $\Rightarrow O(E)$

If we maintain edges in order

- use binary search to insert/find edge $\Rightarrow O(\log E)$
-

Exercise #2: Array-of-edges Representation

28/105

Assuming an array-of-edges representation ...

Write an algorithm to output all edges of the graph


```

show(g):
|   Input graph g
|
|   for all i=0 to g.nE-1 do
|       print g.edges[i]
|   end for


```

Time complexity: $O(E)$


Adjacency Matrix Representation

30/105

Edges represented by a $V \times V$ matrix



A	0	1	2	3
0	0	1	0	1
1	1	0	0	1
2	0	0	0	1
3	1	1	1	0



A	0	1	2	3
0	0	0	0	1
1	1	0	0	1
2	0	0	0	0
3	0	0	1	0

... Adjacency Matrix Representation

31/105

Advantages

- easily implemented as 2-dimensional array
- can represent graphs, digraphs and weighted graphs
 - graphs: symmetric boolean matrix
 - digraphs: non-symmetric boolean matrix
 - weighted: non-symmetric matrix of weight values

Disadvantages:

- if few edges (sparse) \Rightarrow memory-inefficient

... Adjacency Matrix Representation

32/105

Graph initialisation

```

newGraph(V):
|   Input number of nodes V
|   Output new empty graph
|
|   g.nV = V    // #vertices (numbered 0..V-1)
|   g.nE = 0    // #edges
|   allocate memory for g.edges[][]
|   for all i,j=0..V-1 do
|       g.edges[i][j]=0    // false
|   end for
|   return g

```

... Adjacency Matrix Representation

33/105

Edge insertion

```

insertEdge(g,(v,w)):
|   Input graph g, edge (v,w)
|
|   if g.edges[v][w]=0 then    // (v,w) not in graph
|       g.edges[v][w]=1        // set to true
|       g.edges[w][v]=1
|       g.nE=g.nE+1
|   end if

```

... Adjacency Matrix Representation

34/105

Edge removal

```

removeEdge(g,(v,w)):
|   Input graph g, edge (v,w)
|
|   if g.edges[v][w]≠0 then    // (v,w) in graph
|       g.edges[v][w]=0        // set to false
|       g.edges[w][v]=0
|       g.nE=g.nE-1
|   end if

```

Exercise #3: Show Graph

35/105

Assuming an adjacency matrix representation (undirected graph) ...

Write an algorithm to output all edges of the graph (no duplicates!)

... Adjacency Matrix Representation

36/105

```

show(g):
|   Input graph g
|
|   for all i=0 to g.nV-2 do
|       for all j=i+1 to g.nV-1 do
|           if g.edges[i][j] then
|               print i-"j"
|           end if
|       end for
|   end for

```

Time complexity: $O(V^2)$

Exercise #4:

37/105

Analyse storage cost and time complexity of adjacency matrix representation

Storage cost: $O(V^2)$

If the graph is sparse, most storage is wasted.


Cost of operations:

- initialisation: $O(V^2)$ (initialise $V \times V$ matrix)
- insert edge: $O(1)$ (set two cells in matrix)
- delete edge: $O(1)$ (unset two cells in matrix)

... Adjacency Matrix Representation

38/105

A storage optimisation: store only top-right part of matrix.




New storage cost: $V-1$ int ptrs + $V(V-1)/2$ ints (but still $O(V^2)$)

Requires us to always use edges (v,w) such that $v < w$.

Adjacency List Representation

39/105

For each vertex, store linked list of adjacent vertices:



... Adjacency List Representation

40/105

Advantages

- relatively easy to implement in languages like C
- can represent graphs and digraphs
- memory efficient if $E:V$ relatively small

Disadvantages:

- one graph has many possible representations of the lists (unless lists are ordered by same criterion e.g. ascending)

... Adjacency List Representation

41/105

Graph initialisation

```

newGraph(V):
|   Input   number of nodes V
|   Output new empty graph
|
|   g.nV = V    // #vertices (numbered 0..V-1)
|   g.nE = 0    // #edges
|   allocate memory for g.edges[]
|   for all i=0..V-1 do
|       g.edges[i]=NULL    // empty list
|   end for
|   return g

```

... Adjacency List Representation

42/105

Edge insertion:

```

insertEdge(g,(v,w)):
|   Input graph g, edge (v,w)
|
|   insertLL(g.edges[v],w)
|   insertLL(g.edges[w],v)
|   g.nE=g.nE+1

```

... Adjacency List Representation

43/105

Edge removal:

```

removeEdge(g,(v,w)):
|   Input graph g, edge (v,w)
|
|   deleteLL(g.edges[v],w)
|   deleteLL(g.edges[w],v)
|   g.nE=g.nE-1

```

Exercise #5:

44/105

Analyse storage cost and time complexity of adjacency list representation

Storage cost: $O(V+E)$ (V list pointers, total of $2 \cdot E$ list elements)

- the larger of V, E determines the complexity

Cost of operations:

- initialisation: $O(V)$ (initialise V lists)
- insert edge: $O(1)$ (insert one vertex into list)
 - if you don't check for duplicates
- find/delete edge: $O(V)$ (need to find vertex in list)

If vertex lists are sorted

- insert requires search of list $\Rightarrow O(V)$
 - delete always requires a search, regardless of list order
-

Comparison of Graph Representations

45/105

	array of edges	adjacency matrix	adjacency list

space usage	E	V^2	$V+E$
initialise	I	V^2	V
insert edge	I	I	I
find/delete edge	E	I	V

Other operations:

	array of edges	adjacency matrix	adjacency list
disconnected(v)?	E	V	I
isPath(x,y)?	$E \cdot \log V$	V^2	$V+E$
copy graph	E	V^2	E
destroy graph	I	V	E

Graph Abstract Data Type

Graph ADT

47/105

Data:

- set of edges, set of vertices

Operations:

- building: create graph, add edge
- deleting: remove edge, drop whole graph
- scanning: check if graph contains a given edge

Things to note:

- set of vertices is fixed when graph initialised
- we treat vertices as ints, but could be arbitrary Items

... Graph ADT

48/105

Graph ADT interface **graph.h**

```
// graph representation is hidden
typedef struct GraphRep *Graph;

// vertices denoted by integers 0..N-1
typedef int Vertex;

// edges are pairs of vertices (end-points)
typedef struct Edge { Vertex v; Vertex w; } Edge;

// operations on graphs
Graph newGraph(int V);           // new graph with V vertices
int  numOfVertices(Graph);       // get number of vertices in a graph
void insertEdge(Graph, Edge);
void removeEdge(Graph, Edge);
bool adjacent(Graph, Vertex, Vertex); /* is there an edge
```

```
void showGraph(Graph);
void freeGraph(Graph);
```


```
between two vertices */
// print all edges in a graph
```

Exercise #6: Graph ADT Client

49/105

Write a program that uses the graph ADT to

- build the graph depicted below
- print all the nodes that are incident to vertex 1 in ascending order



```
#include <stdio.h>
#include "Graph.h"

#define NODES 4
#define NODE_OF_INTEREST 1

int main(void) {
    Graph g = newGraph(NODES);

    Edge e;
    e.v = 0; e.w = 1; insertEdge(g,e);
    e.v = 0; e.w = 3; insertEdge(g,e);
    e.v = 1; e.w = 3; insertEdge(g,e);
    e.v = 3; e.w = 2; insertEdge(g,e);

    int v;
    for (v = 0; v < NODES; v++) {
        if (adjacent(g, v, NODE_OF_INTEREST))
            printf("%d\n", v);
    }



    freeGraph(g);
    return 0;
}
```

Graph ADT (Array of Edges)

51/105

Implementation of GraphRep (array-of-edges representation)

```
typedef struct GraphRep {
    Edge *edges; // array of edges
    int nV;      // #vertices (numbered 0..nV-1)
    int nE;      // #edges
    int n;       // size of edge array
} GraphRep;
```




Graph ADT (Adjacency Matrix)

52/105

Implementation of GraphRep (adjacency-matrix representation)

```
typedef struct GraphRep {
    int **edges; // adjacency matrix
    int  nV;     // #vertices
    int  nE;     // #edges
} GraphRep;
```



... Graph ADT (Adjacency Matrix)

53/105

Implementation of graph initialisation (adjacency-matrix representation)

```
Graph newGraph(int V) {
    assert(V >= 0);
    int i;

    Graph g = malloc(sizeof(GraphRep));    assert(g != NULL);
    g->nV = V;  g->nE = 0;

    // allocate memory for each row
    g->edges = malloc(V * sizeof(int *));    assert(g->edges != NULL);
    // allocate memory for each column and initialise with 0
    for (i = 0; i < V; i++) {
        g->edges[i] = calloc(V, sizeof(int)); assert(g->edges[i] != NULL);
    }
    return g;
}
```

standard library function **calloc(size_t nelems, size_t nbytes)**

- allocates a memory block of size nelems*nbytes
- and sets all bytes in that block to zero

... Graph ADT (Adjacency Matrix)

54/105

Implementation of edge insertion/removal (adjacency-matrix representation)

```
// check if vertex is valid in a graph
bool validV(Graph g, Vertex v) {
    return (g != NULL && v >= 0 && v < g->nV);
}

void insertEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (!g->edges[e.v][e.w]) { // edge e not in graph
        g->edges[e.v][e.w] = 1;
        g->edges[e.w][e.v] = 1;
        g->nE++;
    }
}
```

```
void removeEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (g->edges[e.v][e.w]) {    // edge e in graph
        g->edges[e.v][e.w] = 0;
        g->edges[e.w][e.v] = 0;
        g->nE--;
    }
}
```

Exercise #7: Checking Neighbours (i)

55/105

Assuming an adjacency-matrix representation ...

Implement a function to check whether two vertices are directly connected by an edge

```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```

```
bool adjacent(Graph g, Vertex x, Vertex y) {
    assert(g != NULL && validV(g,x) && validV(g,y));

    return (g->edges[x][y] != 0);
}
```


Graph ADT (Adjacency List)

57/105

Implementation of GraphRep (adjacency-list representation)

```
typedef struct GraphRep {
    Node **edges; // array of lists
    int    nV;    // #vertices
    int    nE;    // #edges
} GraphRep;

typedef struct Node {
    Vertex    v;
    struct Node *next;
} Node;
```



Exercise #8: Checking Neighbours (ii)

58/105

Assuming an adjacency list representation ...

Implement a function to check whether two vertices are directly connected by an edge

```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```

```
bool adjacent(Graph g, Vertex x, Vertex y) {
    assert(g != NULL && validV(g,x));
```



```
    return inLL(g->edges[x], y);
}
```

inLL() checks if linked list contains an element

Problems on Graphs

Problems on Graphs

61/105

What kind of problems do we want to solve on/via graphs?

- is the graph fully-connected (# connected components is <2)?
 - can we remove an edge and keep it fully-connected?
 - is one vertex reachable starting from some other vertex?
 - which vertices are reachable from v ? (transitive closure)
 - is there a cycle that passes through all vertices? (circuit)
 - what is the cheapest cost path from v to w ?
 - is there a tree that links all vertices? (spanning tree)
 - what is the minimum spanning tree?
 - ...
 - can a graph be drawn in a plane with no crossing edges? (planar graphs)
 - are two graphs "equivalent"? (isomorphism)
-

Graph Algorithms

62/105

In this course we examine algorithms for

- graph traversal (simple graphs)
 - reachability (directed graphs)
 - minimum spanning trees (weighted graphs)
 - shortest path (weighted graphs)
-

Graph Traversal

Finding a Path

64/105

Questions on paths:

- is there a path between two given vertices ($src, dest$)?
- what is the sequence of vertices from src to $dest$?

Approach to solving problem:


- examine vertices adjacent to src
- if any of them is $dest$, then done
- otherwise examine vertices two edges from src (neighbours of neighbours)
- repeat looking further and further from src

Two strategies for graph traversal/search: *depth-first*, *breadth-first*

- DFS follows one path to completion before considering others
 - BFS "fans-out" from the starting vertex ("spreading" subgraph)
-

... Finding a Path

65/105

Comparison of BFS/DFS search for checking if there is a path from a to h ...

Both approaches ignore some edges by remembering previously visited vertices.

Depth-first Search

66/105

Depth-first search can be described recursively as

depthFirst(G, v):

1. mark v as visited
2. for each $(v, w) \in \text{edges}(G)$ do
 - if w has not been visited then
 - depthFirst(w)**

The recursion induces *backtracking*

... Depth-first Search

67/105

Recursive DFS path checking

```


hasPath( $G, \text{src}, \text{dest}$ ):
|   Input  graph  $G$ , vertices  $\text{src}, \text{dest}$ 
|   Output true if there is a path from  $\text{src}$  to  $\text{dest}$  in  $G$ ,
|           false otherwise
|
|   mark all vertices in  $G$  as unvisited
|   return dfsPathCheck( $G, \text{src}, \text{dest}$ )

dfsPathCheck( $G, v, \text{dest}$ ):
|   mark  $v$  as visited
|   if  $v = \text{dest}$  then           // found dest
|       return true
|   else
|       for all  $(v, w) \in \text{edges}(G)$  do
|           if  $w$  has not been visited then
|               return dfsPathCheck( $G, w, \text{dest}$ ) // found path via  $w$  to dest
|           end if
|       end for
|   end if
|   return false           // no path from  $v$  to dest
  
```

Exercise #9: Depth-first Traversal (i)

68/105

Trace the execution of `dfsPathCheck($G, 0, 5$)` on:



Consider neighbours in ascending order

Answer:

0 - 1 - 2 - 3 - 4 - 5

... Depth-first Search

70/105

Cost analysis:

- all vertices marked as unvisited, each vertex visited at most once \Rightarrow cost = $O(V)$
- visit all edges incident on visited vertices \Rightarrow cost = $O(E)$
 - assuming an adjacency list representation

Time complexity of DFS: $O(V+E)$ (adjacency list representation)

- the larger of V, E determines the complexity
-

... Depth-first Search

71/105

Note how different graph data structures affect cost:

- array-of-edges representation
 - visit all edges incident on visited vertices \Rightarrow cost = $O(V \cdot E)$
 - cost of DFS: $O(V \cdot E)$
- adjacency-matrix representation
 - visit all edges incident on visited vertices \Rightarrow cost = $O(V^2)$
 - cost of DFS: $O(V^2)$

In case of Adjacency List:

For *dense graphs* ... $E \cong V^2 \Rightarrow O(V+E) = O(V^2)$

For *sparse graphs* ... $E \cong V \Rightarrow O(V+E) = O(E)$

... Depth-first Search

72/105

Knowing whether a path exists can be useful

Knowing what the path is even more useful

\Rightarrow record the previously visited node as we search through the graph (so that we can then trace path through graph)

Make use of global variable:

- `visited[]` ... array to store previously visited node, for each node being visited


visited[] // store previously visited node, for each vertex 0..nV-1

```
findPath(G,src,dest):
  Input graph G, vertices src,dest

  for all vertices v∈G do
    visited[v]=-1
  end for
  visited[src]=src // starting node of the path
  if dfsPathCheck(G,src,dest) then // show path in dest..src order
    v=dest
    while v≠src do
      print v "-"
      v=visited[v]
    end while
    print src
  end if

dfsPathCheck(G,v,dest):
  if v=dest then // found edge from v to dest
    return true
  else
    for all (v,w)∈Edges(G) do
      if visited[w]=-1 then
        visited[w]=v
        if dfsPathCheck(G,w,dest) then
          return true // found path via w to dest
        end if
      end if
    end for
  end if
  return false // no path from v to dest
```

Show the DFS order in which we visit vertices in this graph when searching for a path from 0 to 6:



Consider neighbours in ascending order

0	0	3	5	3	1	5	4	7	8
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

Path: 6-5-1-0

DFS can also be described non-recursively (via a *stack*):

hasPath(G,src,dest):

Input graph G, vertices src,dest

Output true if there is a path from src to dest in G,
false otherwise

mark all vertices in G as unvisited

push src onto new stack s

found=false

while not found and s is not empty **do**

 pop v from s

 mark v as visited

if v=dest **then**

 found=true

else

for each (v,w)∈Edges(G) such that w has not been visited

 push w onto s

end for

end if

end while

return found


Uses standard stack operations (push, pop, check if empty)

Time complexity is the same: $O(V+E)$

Exercise #11: Depth-first Traversal (iii)

77/105

Show how the stack evolves when executing findPathDFS(g,0,5) on:



Push neighbours in *descending* order ... so they get popped in ascending order

						4	5
				3	5	5	5
	1	2	4	4	4	4	4
	4	4	4	4	4	4	4
(empty)	→ 0	→ 5	→ 5	→ 5	→ 5	→ 5	→ 5

Breadth-first Search

79/105

Basic approach to breadth-first search (BFS):

- visit and mark current vertex
- visit all neighbours of current vertex
- then consider neighbours of neighbours

Notes:

- tricky to describe recursively
- a minor variation on non-recursive DFS search works
⇒ switch the *stack* for a *queue*

... Breadth-first Search

80/105

BFS algorithm (records visiting order, marks vertices as visited when put *on* queue):

visited[] // array of visiting orders, indexed by vertex 0..nV-1

```

findPathBFS(G,src,dest):
|   Input  graph G, vertices src,dest
|
|   for all vertices v∈G do
|       visited[v]=-1
|   end for
|   enqueue src into new queue q
|   visited[src]=src
|   found=false
|   while not found and q is not empty do
|       dequeue v from q
|       if v=dest then
|           found=true
|       else
|           for each (v,w)∈Edges(G) such that visited[w]=-1 do
|               enqueue w into q
|               visited[w]=v
|           end for
|       end if
|   end while
|   if found then
|       display path in dest..src order
|   end if


```

Uses standard queue operations (enqueue, dequeue, check if empty)

Exercise #12: Breadth-first Traversal

81/105

Show the BFS order in which we visit vertices in this graph when searching for a path from 0 to 6:



Consider neighbours in ascending order

0	0	0	2	5	0	5	5	3	-1
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

... Breadth-first Search

83/105

Time complexity of BFS: $O(V+E)$ (adjacency list representation, same as DFS)

BFS finds a "shortest" path

- based on minimum # edges between *src* and *dest*.
- stops with first-found path, if there are multiple ones

In many applications, edges are weighted and we want path

- based on minimum sum-of-weights along path *src* .. *dest*


We discuss weighted/directed graphs later.

Other DFS Examples

84/105

Other problems to solve via DFS graph search

- checking for the existence of a cycle
- determining which connected component each vertex is in



Exercise #13: Buggy Cycle Check

85/105

A graph has a *cycle* if

- it has a path of length > 1
- with start vertex *src* = end vertex *dest*
- and without using any edge more than once

We are not required to give the path, just indicate its presence.

The following DFS cycle check has two bugs. Find them.

```
hasCycle(G):
|   Input   graph G
|   Output true if G has a cycle, false otherwise
|
|   choose any vertex  $v \in G$ 
|   return dfsCycleCheck(G,v)
|
dfsCycleCheck(G,v):
|   mark v as visited
|   for each (v,w)  $\in$  Edges(G) do
|   |   if w has been visited then // found cycle
|   |   |   return true
|   |   else if dfsCycleCheck(G,w) then
|   |   |   return true
```

```
| end for
| return false           // no cycle at v
```

1. Only one connected component is checked.
2. The loop

for each $(v,w) \in \text{Edges}(G)$ **do**

should exclude the neighbour of v from which you just came, so as to prevent a single edge $w-v$ from being classified as a cycle.

Computing Connected Components

87/105

Problems:

- how many connected subgraphs are there?
- are two vertices in the same connected subgraph?

Both of the above can be solved if we can

- build an array, one element for each vertex V
 - indicating which connected component V is in
 - `componentOf[]` ... array $[0..nV-1]$ of component IDs
-

... Computing Connected Components

88/105

Algorithm to assign vertices to connected components:

```
components(G):
| Input graph G
|
| for all vertices  $v \in G$  do
|   componentOf[v] = -1
| end for
| compID = 0
| for all vertices  $v \in G$  do
|   if componentOf[v] == -1 then
|     dfsComponents(G, v, compID)
|     compID = compID + 1
|   end if
| end for


dfsComponents(G, v, id):
| componentOf[v] = id
| for all vertices  $w$  adjacent to  $v$  do
|   if componentOf[w] == -1 then
|     dfsComponents(G, w, id)
|   end if
| end for
```

Exercise #14: Connected components

89/105

Trace the execution of the algorithm

1. on the graph shown below
2. on the same graph but with the dotted edges added



Consider neighbours in ascending order

1.

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
-1	-1	-1	-1	-1	-1	-1	-1
0	-1	-1	-1	-1	-1	-1	-1
0	-1	0	-1	-1	-1	-1	-1
0	0	0	-1	-1	-1	-1	-1
0	0	0	1	-1	-1	-1	-1
...							
0	0	0	1	1	2	2	2

2.

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
-1	-1	-1	-1	-1	-1	-1	-1
0	-1	-1	-1	-1	-1	-1	-1
0	0	-1	-1	-1	-1	-1	-1
0	0	0	-1	-1	-1	-1	-1
...							
0	0	0	0	0	1	1	1

Hamiltonian and Euler Paths

Hamiltonian Path and Circuit

92/105

Hamiltonian path problem:

- find a path connecting two vertices v, w in graph G
- such that the path includes each *vertex* in G exactly once

If $v = w$, then we have a *Hamiltonian circuit*

Simple to state, but difficult to solve (*NP*-complete)

Many real-world applications require you to visit all vertices of a graph:


- Travelling salesman
- Bus routes
- ...

Problem named after Irish mathematician, physicist and astronomer Sir William Rowan Hamilton (1805 — 1865)

... Hamiltonian Path and Circuit

93/105

Graph and two possible Hamiltonian paths:



... Hamiltonian Path and Circuit

94/105

Approach:

- generate all possible simple paths (using e.g. DFS)
- keep a counter of vertices visited in current path
- stop when find a path containing V vertices

Can be expressed via a recursive DFS algorithm

- similar to simple path finding approach, except
 - keeps track of path length; succeeds if length = $v-1$ (length = v for circuit)
 - resets "visited" marker after unsuccessful path

... Hamiltonian Path and Circuit

95/105

Algorithm for finding Hamiltonian path:

visited[] // array [0.. $nV-1$] to keep track of visited vertices

hasHamiltonianPath(G,src,dest):

```

|   for all vertices  $v \in G$  do
|       visited[v]=false
|   end for
|   return hamiltonR(G,src,dest,#vertices(G)-1)

```

hamiltonR(G,v,dest,d):

```


|   Input G      graph
|           v      current vertex considered
|           dest  destination vertex
|           d      distance "remaining" until path found
|
|   if v=dest then
|       if d=0 then return true else return false
|   else
|       mark v as visited
|       for each unvisited neighbour w of v in G do
|           if hamiltonR(G,w,dest,d-1) then
|               return true
|           end if
|       end for
|   end if
|   mark v as unvisited           // reset visited mark
|                               // no hamiltonian path from v in this consideration
|   return false

```

Exercise #15: Hamiltonian Path

96/105

Trace the execution of the algorithm when searching for a Hamiltonian path from 1 to 6:



Consider neighbours in ascending order

1-0-2-3-4-5-6	$d \neq 0$
1-0-2-3-4-5-7-8-9	no unvisited neighbour
1-0-2-3-4-5-7-9-8	no unvisited neighbour
1-0-2-3-4-7-5-6	$d \neq 0$
1-0-2-3-4-7-8-9	no unvisited neighbour
1-0-2-3-4-7-9-8	no unvisited neighbour
1-0-2-3-4-8-7-5-6	$d \neq 0$
1-0-2-3-4-8-7-9	no unvisited neighbour
1-0-2-3-4-8-9-7-5-6	✓


Repeat on your own with $src=0$ and $dest=6$

... Hamiltonian Path and Circuit

98/105

Analysis: worst case requires $(V-1)!$ paths to be examined

Consider a graph with isolated vertex and the rest fully-connected



Checking $\text{hasHamiltonianPath}(g, x, 0)$ for any x

- requires us to consider every possible path
- e.g 1-2-3-4, 1-2-4-3, 1-3-2-4, 1-3-4-2, 1-4-2-3, ...
- starting from any x , there are $3!$ paths $\Rightarrow 4!$ total paths
- there is no path of length 5 in these $(V-1)!$ possibilities

There is no known simpler algorithm for this task

Note, however, that the above case could be solved in constant time if we had a fast check for 0 and x being in the same connected component


Euler Path and Circuit

99/105


Euler path problem:

- find a path connecting two vertices v, w in graph G
- such that the path includes each *edge* exactly once
(note: the path does not have to be simple \Rightarrow can visit vertices more than once)

If $v = w$, then we have an *Euler circuit*



Euler Path: 4-3-1-5-2-1



Euler Circuit: 1-2-5-4-3-1

Many real-world applications require you to visit all edges of a graph:

- Postman
- Garbage pickup
- ...

Problem named after Swiss mathematician, physicist, astronomer, logician and engineer Leonhard Euler (1707 — 1783)

... Euler Path and Circuit

100/105

One possible "brute-force" approach:

- check for each path if it's an Euler path
- would result in factorial time performance

Can develop a better algorithm by exploiting:



Theorem. A graph has an *Euler circuit* if and only if it is connected and all vertices have even degree

Theorem. A graph has a *non-circuitous Euler path* if and only if it is connected and exactly two vertices have odd degree

Exercise #16: Euler Paths and Circuits

101/105

Which of these two graphs have an Euler path? an Euler circuit?



No Euler circuit

Only the second graph has an Euler path, e.g. 2-0-1-3-2-4-5

... Euler Path and Circuit

103/105

Assume the existence of $\text{degree}(g, v)$ (degree of a vertex)

Algorithm to check whether a graph has an Euler path:

```

hasEulerPath(G, src, dest):
|   Input  graph G, vertices src, dest
|   Output true if G has Euler path from src to dest
|           false otherwise
|
|   if degree(G, src) is even or degree(G, dest) is even then
|       return false
|   end if
|   for all vertices  $v \in G$  do
|       if  $v \neq \text{src}$  and  $v \neq \text{dest}$  and degree(G, v) is odd then
|           return false
|       end if
|   end for
|   return true

```

The basic idea: ensure that all vertices have even degree except for the src and dest, which have odd degrees. All non-source/destination vertices with odd degrees must be connected in pairs to form an even degree.

... Euler Path and Circuit

104/105

Analysis of hasEulerPath algorithm:

- assume that connectivity is already checked
- assume that degree is available via $O(1)$ lookup
- single loop over all vertices $\Rightarrow O(V)$

If degree requires iteration over vertices

- cost to compute degree of a single vertex is $O(V)$
- overall cost is $O(V^2)$

\Rightarrow problem tractable, even for large graphs (unlike Hamiltonian path problem)

For the keen, a linear-time (in the number of edges, E) algorithm to compute an Euler path is described in [Sedgewick] Ch.17.7.

Summary

105/105

- Graph terminology
 - vertices, edges, vertex degree, connected graph, tree
 - path, cycle, clique, spanning tree, spanning forest
- Graph representations
 - array of edges
 - adjacency matrix
 - adjacency lists
- Graph traversal
 - depth-first search (DFS)
 - breadth-first search (BFS)
 - cycle check, connected components
 - Hamiltonian paths/circuits, Euler paths/circuits

- Suggested reading (Sedgewick):
 - graph representations ... Ch. 17.1-17.5
 - Hamiltonian/Euler paths ... Ch. 17.7
 - graph search ... Ch. 18.1-18.3, 18.7
-

Produced: 5 Mar 2024